

Name: ANUJA PATIL
Student ID: 862123752

1) Change exit implementation to maintain exit status.

Approach: I have added another implementation rather than modifying the existing exit(). The name of newly added exit is **exitUDef()**. To make the required changes, I changed some files. Following are the files changed and a brief summary of the changes done in each:

a) syscall.h

Defined a system call and associated it with a number:

#define SYS_exitUDef 22

b) syscall.c

Added the function prototype for exitUDef() – **extern int sys_exitUDef(void);**

Added a pointer to the system call in the array of function pointers:

```
static int(*syscalls[]) (void) = {  
    .  
    .  
    .  
    [SYS_exitUDef] sys_exitUDef,  
}
```

c) usys.S

In order for a user program to call the system call, an interface was added:

SYS_CALL(exitUDef)

d) sysproc.c

Added the handler for exitUDef system call to maintain exit status:

```
int sys_exitUDef(void) {  
    int exitStatus;  
    if(argint(0, &exitStatus) < 0) return -1;  
    return exitUDef(exitStatus);  
}
```

e) proc.h

Added a field called exitStatus in the proc struct which stores the state of a process:

```
struct proc {  
    .  
    .  
    int exitStatus;  
}
```

f) proc.c

This contains the actual implementation of the `exitUDef()`. This is same as original `exit()` except that in `exitUDef()`, we also store status received as an argument in the current process's `exitStatus` state:

```
curproc->exitStatus = status;
```

where `status` is the argument passed to `exitUDef()`.

g) defs.h

Added forward declaration of system call `exitUDef()` in “`proc.c`” section:

```
int exitUDef(int);
```

h) user.h

Added user level definition of `exitUDef()` under system calls:

```
int exitUDef(int);
```

2) Update the `wait` system call signature to `int wait(int *status)`.

Approach: Updated the existing `wait()` to contain the pointer to integer argument ‘`status`’.

Original `wait()` returns `pid` of the process that current process was waiting on. If we want to return another value, in our case `exit status` of the process too, then we pass the address of a variable to `wait()` so that the `exit status` of the process can be stored in the passed argument. To make the required changes, I changed some files. Following are the files changed and a brief summary of the changes done in each:

a) sysproc.c

Added a field ‘`status`’ which is an integer pointer. Called `argptr()` to pass the address of ‘`status`’:

```
int sys_wait(void) {  
    int* status;  
    argptr(0, (void*)&status, sizeof(status));  
    return wait(status);  
}
```

b) proc.c

Modified the `wait()` method to contain the pointer to integer argument ‘`status`’. If ‘`status`’ is not null (which indicates that the user also wants to return the `exit status` of the process), we store the process's `exitStatus` in ‘`status`’ field:

```
int wait(int* status) {  
    .  
    .  
    if(p->state == ZOMBIE) {  
        .  
        .  
        p->state = UNUSED;  
        if(status) *status = p->exitStatus;  
        p->exitStatus = 0; //reset to 0  
        .  
    }
```

c) defs.h

Modified forward declaration of system call wait() in "proc.c" section:

```
int wait(int*);
```

d) user.h

Modified user level definition of wait() under system calls:

```
int wait(int*);
```

3) Add a waitpid system call: int waitpid(int pid, int *status, int options).

Approach: Added implementation for **waitpid(int pid, int* status, int options)**. To make the required changes, I changed some files. Following are the files changed and a brief summary of the changes done in each:

a) syscall.h

Defined a system call and associated it with a number:

```
#define SYS_waitpid 23
```

b) syscall.c

Added the function prototype for waitpid() – **extern int sys_waitpid(void);**

Added a pointer to the system call in the array of function pointers:

```
static int(*syscalls[]) (void) = {  
    .  
    .  
    .  
    [SYS_waitpid]  sys_waitpid,  
}
```

c) usys.S

In order for a user program to call the system call, an interface was added:

```
SYS_CALL(waitpid)
```

d) sysproc.c

Added the handler for waitpid system call to maintain exit status:

```
int sys_waitpid (void) {  
    int pid, options = 0;  
    int* status;  
    if(argint(0, &pid) < 0) return -1;  
    if(argptr(1, (void*)&status, sizeof(status)) < 0) return -1;  
    return waitpid(pid, status, options);  
}
```

e) defs.h

Added forward declaration of system call waitpid() in "proc.c" section:

```
int waitpid(int, int*, int);
```

f) user.h

Added user level definition of waitpid() under system calls:

```
int waitpid(int, int*, int);
```

g) proc.c

This contains the actual implementation of the waitpid(). This is similar to the modified wait() except that in waitpid(), we are not looping for the current process's children only. We are looping over all the processes in the process table and looking for the process whose pid matches with the one sent as a parameter to the waitpid(). I have renamed 'havekids' to 'processExists':

```
for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {  
    if(p->pid != pid) continue;  
    processExists = 1;  
    if(p->state == ZOMBIE) {  
        .  
        .  
        p->state = UNUSED;  
        if(status) *status = p->exitStatus;  
        p->exitStatus = 0;  
        release(&ptable.lock);  
        return pid;  
    }  
    .  
    .  
}
```

4) Modified the lab1_testfile.c to use exitUDef(x) in place of exit(x) where x is an exit status in the program. I have also changed **wait()** to **wait(0)** in files – forktest.c, init.c, sh.c, stressfs.c, usertests.c. 0 parameter in wait() indicates that we are passing null. In C, 0 means null.

5) BONUS – Implement WNOHANG and create version of CELEBW02 on same page that check of a child process still running.

Approach: lab1_testfile.c has the WNOHANG defined to be 1. I changed some files to implement WNOHANG as required in the assignment. Following are the files changed and a brief summary of the changes done in each:

a) sysproc.c

We need to check if any options argument has been passed to the method in waitpid(). By default, options takes the value 0. For this assignment, according to lab1_testfile, options can take the value 1:

```
int sys_ waitpid (void) {  
    int pid, options = 0;  
    int* status;  
    if(argint(0, &pid) < 0) return -1;  
    if(argptr(1, (void*)&status, sizeof(status)) < 0) return -1;
```

```

    if(argint(2, &options) < 0) return -1;
    return waitpid(pid, status, options);
}

```

b) proc.c

We just want to implement WNOHANG such that lab1_testfile.c runs. In this test file, for the bonus part, we call waitpid(pid, &status, WNOHANG) where WNOHANG is defined to be 1. If the waitpid returns -1, there is an error with wait(); if it returns pid, it means the process with given pid has terminated; if it returns 0, it means the process with given pid is still running and since we do not want to make the current process wait, waitpid() should return 0. In proc.c, we add an 'else-if' block to the 'if':

```

if(p->state == ZOMBIE) {
    .
    .
} else if(options == 1) {
    release(&ptable.lock);
    return 0;
}

```

The control reaches the 'else-if' block when the process with given pid exists, is not in ZOMBIE state and options is 1.

6) Following are the screenshots of the parts of code in the files mentioned above which I have changed:

a) syscall.h

```

// System call numbers
#define SYS_fork 1
#define SYS_exit 2
#define SYS_wait 3
#define SYS_pipe 4
#define SYS_read 5
#define SYS_kill 6
#define SYS_exec 7
#define SYS_fstat 8
#define SYS_chdir 9
#define SYS_dup 10
#define SYS_getpid 11
#define SYS_sbrk 12
#define SYS_sleep 13
#define SYS_uptime 14
#define SYS_open 15
#define SYS_write 16
#define SYS_mknod 17
#define SYS_unlink 18
#define SYS_link 19
#define SYS_mkdir 20
#define SYS_close 21
#define SYS_exitUDef 22
#define SYS_waitpid 23
~
~
"syscall.h" 24L, 532C

```

b) syscall.c

```
apati027@sledge:~/xv6 — ssh a...    ...h apati027@sledge.cs.u...
extern int sys_exitUDef(void);
extern int sys_waitpid(void);
ited with status %d
```

```
apati027@sledge:~/xv6 — ssh a...    ...h apati027@sledge.cs.ucr.edu
child with PID# %d
static int (*syscalls[])(void) = {
[SYS_fork] sys_fork,
[SYS_exit] sys_exit,
[SYS_wait] sys_wait,
[SYS_pipe] sys_pipe,
[SYS_read] sys_read,
[SYS_kill] sys_kill,
[SYS_exec] sys_exec,
[SYS_fstat] sys_fstat,
[SYS_chdir] sys_chdir,
[SYS_dup] sys_dup,
[SYS_getpid] sys_getpid,
[SYS_sbrk] sys_sbrk,
[SYS_sleep] sys_sleep,
[SYS_uptime] sys_uptime,
[SYS_open] sys_open,
[SYS_write] sys_write,
[SYS_mknod] sys_mknod,
[SYS_unlink] sys_unlink,
[SYS_link] sys_link,
[SYS_mkdir] sys_mkdir,
[SYS_close] sys_close,
[SYS_exitUDef] sys_exitUDef,
[SYS_waitpid] sys_waitpid,
};
```

c) usys.S

```
apati027@sledge:~/xv6 — ssh a...    ...h apati027@sledge.cs.ucr.edu
int $T_SYSCALL; \
ret

Lectures    untitled folder
SYSCALL(fork)
SYSCALL(exit)
SYSCALL(wait)
SYSCALL(pipe)
SYSCALL(read)
SYSCALL(write)
SYSCALL(close)
SYSCALL(kill)
SYSCALL(exec)
SYSCALL(open)
SYSCALL(mknod)
SYSCALL(unlink)
SYSCALL(fstat)
SYSCALL(link)
SYSCALL(mkdir)
SYSCALL(chdir)
SYSCALL(dup)
SYSCALL(getpid)
SYSCALL(sbrk)
SYSCALL(sleep)
SYSCALL(uptime)
SYSCALL(exitUDef)
SYSCALL(waitpid)
"usys.S" 33L, 496C
```

d) sysproc.c

```
apati027@sledge:~/xv6 — ssh a...    ...h apati027@sledge.cs.ucr.edu
int
sys_exitUDef(void)
{
    int exitStatus;
    if(argint(0, &exitStatus) < 0) return -1;
    return exitUDef(exitStatus);
}
```

```
// Changed for assignment1
int
sys_wait(void)
{
    int* status;
    argptr(0, (void*)&status, sizeof(status));
    return wait(status);
}
```

```
int
sys_waitpid(void)
{
    int pid, options = 0;
    int* status;
    if(argint(0, &pid) < 0) return -1;
    if(argptr(1, (void*)&status, sizeof(status)) < 0) return -1;
    if(argint(2, &options) < 0) return -1;
    return waitpid(pid, status, options);
}

"sysproc.c" 125L, 2223C
```


e) proc.h

```

enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };

// Per-process state
struct proc {
    uint sz; // Size of process memory (bytes)
    pde_t* pgdir; // Page table
    char *kstack; // Bottom of kernel stack for this process
    enum procstate state; // Process state
    int pid; // Process ID
    struct proc *parent; // Parent process
    struct trapframe *tf; // Trap frame for current syscall
    struct context *context; // user context to run process
    void *chan; // Channel for I/O
    int killed; // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd; // Current directory
    char name[16]; // Process name (debugging)
    int exitStatus; // Exit status (added for assignment1)
};

```

f) proc.c

```

//Added for assignment1.
//Defining exitUDef where we store exit status (passed as a parameter)
//into process' state. Rest everything is same as that of exit()
int
exitUDef(int status)
{
    struct proc *curproc = myproc();
    struct proc *p;
    int fd;

    if(curproc == initproc)
        panic("init exiting");

    // Close all open files.
    for(fd = 0; fd < NOFILE; fd++){
        if(curproc->ofile[fd]){
            fileclose(curproc->ofile[fd]);
            curproc->ofile[fd] = 0;
        }
    }

    begin_op();
    iput(curproc->cwd);
    end_op();
    curproc->cwd = 0;
}

```

```
apati027@sledge:~/xv6 — ssh a...  ...h apati027@sledge.cs.ucr.edu  ~ — -bash

begin_op();
input(curproc->cwd);
end_op();
curproc->cwd = 0;

acquire(&ptable.lock);

// Parent might be sleeping in wait().
wakeup1(curproc->parent);

// Pass abandoned children to init.
for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
    if(p->parent == curproc){
        p->parent = initproc;
        if(p->state == ZOMBIE)
            wakeup1(initproc);
    }
}

// Jump into the scheduler, never to return.
curproc->state = ZOMBIE;
curproc->exitstatus = status;
sched();
panic("zombie exit");
}
```

```
apati027@sledge:~/xv6 — ssh a...  ...h apati027@sledge.cs.ucr.edu  ~ — -bash  ~/Desktop — -bash

// Added for assignment1.
// This method waits for a process (not necessary a child process) with a pid
// that equals to the one provided by the pid argument. The return value must
// be the process id of the process that was terminated or -1 if this process
// does not exist or if an unexpected error occurred.
// This method is similar to the wait() defined above.
int
waitpid(int pid, int* status, int options)
{
    struct proc *p, *curproc = myproc();
    int processExists; //similar to havekids in wait()
    acquire(&ptable.lock);

    //Loops continuously till the process with given pid is terminated
    for(;;){
        processExists = 0;

        //Scan through the process table looking for exited processes. Terminated
        //processes will be in ZOMBIE state.
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            // If the process pid does not match the given pid, no need to continue
            // with this process.
            if(p->pid != pid) continue;

            processExists = 1;
        }
    }
}
```

```
apati027@sledge:~/xv6 — ssh a...  ...h apati027@sledge.cs.ucr.edu  ~ — -bash  ~/Desktop — -bash

processExists = 1;
if(p->state == ZOMBIE) {
    //Found the process with the given pid that has exited.
    kfree(p->kstack);
    p->kstack = 0;
    freevm(p->pgdir);
    p->pid = 0;
    p->parent = 0;
    p->name[0] = 0;
    p->killed = 0;
    p->state = UNUSED;
    if(status) *status = p->exitStatus;
    p->exitStatus = 0;
    release(&ptable.lock);
    return pid;
} else if(options == 1) { //if options is passed by the user.

    //the process with the given pid is still running, so we
    //don't block the current process, just release the lock on
    //ptable and return 0.
    release(&ptable.lock);
    return 0;
}

// No point waiting if the process with given pid does not exist
```

```
// No point waiting if the process with given pid does not exist
// or the current process is killed.
if(!processExists || curproc->killed) {
    release(&ptable.lock);
    return -1;
}

// Wait for the process with the given pid to exit.
sleep(curproc, &ptable.lock);
}
```

g) defs.h

```
apati027@sledge:~/xv6 — ssh a...    ...h apati027@sledge.cs.ucr.edu    ~ — -bash    ~/Des

//PAGEBREAK: 16
// proc.c
int      cpuid(void);
void     exit(void);
int      fork(void);
int      growproc(int);
int      kill(int);
struct cpu* mycpu(void);
struct proc* myproc();
void     pinit(void);
void     procdump(void);
void     scheduler(void) __attribute__((noreturn));
void     sched(void);
void     setproc(struct proc*);
void     sleep(void*, struct spinlock*);
void     userinit(void);
//int     wait(void);
int      wait(int*); //Added for assignment1
void     wakeup(void*);
void     yield(void);
int      exitUDef(int); //Added for assignment1
int      waitpid(int, int*, int); //Added for assignment1
```

h) user.h

```
apati027@sledge:~/xv6 — ssh a...  ...h apati027@sledge.cs.ucr.edu  ~ — -bash
// system calls
int fork(void);
int exit(void) __attribute__((noreturn));
//int wait(void);
int wait(int*); //Added for assignment1
int pipe(int*);
int write(int, const void*, int);
int read(int, void*, int);
int close(int);
int kill(int);
int exec(char*, char**);
int open(const char*, int);
int mknod(const char*, mode_t, ushort);
int unlink(const char*);
int fstat(int fd, struct stat*);
int link(const char*, const char*);
int mkdir(const char*);
int chdir(const char*);
int dup(int);
int getpid(void);
char* sbrk(int);
int sleep(int);
int uptime(void);
int exitUDef(int);
int waitpid(int, int*, int); //Added for assignment1

"user.h" 42L, 1036C
```

7) To make and run the lab1_testfile.c: (We are inside xv6 folder)

> make clean

> make qemu-nox

This will open qemu console. Then run:
\$ lab1_testfile 1

```
apati027@sledge:~/xv6 — ssh apati027@sledge.cs.ucr.edu ~ — apati027

SeaBIOS (version 1.11.0-2.el7)

iPXE (http://ipxe.org) 00:03.0 C980 PCI2.10 PnP PMM+1FF94780+1FED4780 C980

Booting from Hard Disk..xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ lab1_testfile 1

This program tests the correctness of your lab#1

Parts a & b) testing exit(int status) and wait(int* status):

This is child with PID# 4 and I will exit with status 0

This is the parent: child with PID# 4 has exited with status 0

This is child with PID# 5 and I will exit with status -1

This is the parent: child with PID# 5 has exited with status -1
```

\$ lab1_testfile 2

```
apati027@sledge:~/xv6 — ssh apati027@sledge.cs.ucr.edu

$ lab1_testfile 2

This program tests the correctness of your lab#1

Part c) testing waitpid(int pid, int* status, int options):

The is child with PID# 7 and I will exit with status 11
The is child with PID# 9 and I will exit with status 13
The is child with PID# 8 and I will exit with status 12
The is child with PID# 10 and I will exit with status 14

The is child with PID# 11 and I will exit with status 15
This is the parent: Now waiting for child with PID# 10
This is the parent: Child# 10 has exited with status 14
This is the parent: Now waiting for child with PID# 8
This is the parent: Child# 8 has exited with status 12
This is the parent: Now waiting for child with PID# 9
This is the parent: Child# 9 has exited with status 13
This is the parent: Now waiting for child with PID# 7
This is the parent: Child# 7 has exited with status 11
This is the parent: Now waiting for child with PID# 11
```

\$ lab1_testfile 3

apati027@sledge:~/xv6 — ssh apati027@sledge.cs.ucr.edu

```
The is child with PID# 9 and I will exit with status 13
The is child with PID# 8 and I will exit with status 12
The is child with PID# 10 and I will exit with status 14

The is child with PID# 11 and I will exit with status 15
This is the parent: Now waiting for child with PID# 10

This is the parent: Child# 10 has exited with status 14

This is the parent: Now waiting for child with PID# 8

This is the parent: Child# 8 has exited with status 12

This is the parent: Now waiting for child with PID# 9

This is the parent: Child# 9 has exited with status 13

This is the parent: Now waiting for child with PID# 7

This is the parent: Child# 7 has exited with status 11

This is the parent: Now waiting for child with PID# 11

This is the parent: Child# 11 has exited with status 15
[$ lab1_testfile 3

This program tests the correctness of your lab#1

Part e) the waitpid option WNOHANG, test program CELEBW02
child is still running
child is still running
child is still running
child is still running
child exited with status of 1
$ █
```

Summary for lab1:

The lab1 was about learning how to create a system call, how the parameters are passed from user space to kernel space and how are parameters returned from kernel space to user space. The main files to change when creating a system call are – syscall.h, syscall.c, sysproc.c, usys.S, user.h, defs.h and proc.c.

syscall.h: mapping from system call name to system call number.

syscall.c: contains pointers to the system calls. We also add a function prototype.

sysproc.c: contains the implementations of process related system calls.

usys.S: contains a list of system calls exported by the kernel.

user.h: contains the system call definitions in xv6.

defs.h: contains forward declaration of system calls.

proc.c: contains the implementations of the system calls, scheduler, exit, wait etc.